

Breast_cancer_prediction

December 13, 2019

1 Breast Cancer Prediction with Machine Learning

1.0.1 1. Motivation

Breast cancer is the most common type of cancer in women around the world and the second highest in terms of mortality rates, and is normally detected when an abnormal lump or a tiny speck of calcium are found (from self-examination or x-ray). Early detection of this kind of cancer can greatly improve survival chances by promoting clinical treatment to patients as soon as possible. That said, the main purpose of this post is to predict whether a lump is benign or malignant through data analysis.

All information and data related to this problem can be found here: <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data/>

1.0.2 2. Data Information

Ten real-valued features are computed for each cell nucleus.

radius (mean of distances from center to points on the perimeter)

texture (standard deviation of gray-scale values)

perimeter

area

smoothness (local variation in radius lengths)

compactness ($\text{perimeter}^2 / \text{area} - 1.0$)

concavity (severity of concave portions of the contour)

concave points (number of concave portions of the contour)

symmetry

fractal dimension ("coastline approximation" - 1)

That said, our dataset consists of 30 features, the mean, standard error and "worst" or largest (mean of the three largest values) of these ten features.

1.0.3 3. Dependences

Here we can find the libraries we will use in order to develop a solution for this problem. **numpy**|**pandas**: Will help us treat the data. **matplotlib**|**seaborn**: Will help us plot the information so we can visualize it in different ways and have a better understanding of it. **sklearn**: Will provide all necessary tools to train our models and test them afterwards. **prettytable**: Will allow us to plot simple ascii tables in order to show results more aesthetically.

```
In [1]: import warnings
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import f1_score, confusion_matrix
from sklearn.ensemble import ExtraTreesClassifier
from sklearn import svm
from sklearn.model_selection import GridSearchCV
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, recall_score
from sklearn.feature_selection import RFECV
from prettytable import PrettyTable
warnings.filterwarnings('ignore')
```

1.0.4 4. Data Cleaning

On this section we will check for undesired values such as NaN and get our dataset prepared so we can train our models later on with proper data. We will also be looking for object data types so we can transform them into numerical values.

```
In [2]: #We load the .csv and visualize the first ten rows of it, here we can also see the col.
df = pd.read_csv("data.csv")
df.head(10)
```

```
Out[2]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	
5	843786	M	12.45	15.70	82.57	477.1	
6	844359	M	18.25	19.98	119.60	1040.0	
7	84458202	M	13.71	20.83	90.20	577.9	
8	844981	M	13.00	21.82	87.50	519.8	
9	84501001	M	12.46	24.04	83.97	475.9	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.30010	0.14710	
1	0.08474	0.07864	0.08690	0.07017	
2	0.10960	0.15990	0.19740	0.12790	
3	0.14250	0.28390	0.24140	0.10520	
4	0.10030	0.13280	0.19800	0.10430	
5	0.12780	0.17000	0.15780	0.08089	

6	0.09463	0.10900	0.11270	0.07400
7	0.11890	0.16450	0.09366	0.05985
8	0.12730	0.19320	0.18590	0.09353
9	0.11860	0.23960	0.22730	0.08543

	...	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	...	17.33	184.60	2019.0	0.1622	
1	...	23.41	158.80	1956.0	0.1238	
2	...	25.53	152.50	1709.0	0.1444	
3	...	26.50	98.87	567.7	0.2098	
4	...	16.67	152.20	1575.0	0.1374	
5	...	23.75	103.40	741.6	0.1791	
6	...	27.66	153.20	1606.0	0.1442	
7	...	28.14	110.60	897.0	0.1654	
8	...	30.73	106.20	739.3	0.1703	
9	...	40.68	97.65	711.4	0.1853	

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
0	0.6656	0.7119	0.2654	0.4601	
1	0.1866	0.2416	0.1860	0.2750	
2	0.4245	0.4504	0.2430	0.3613	
3	0.8663	0.6869	0.2575	0.6638	
4	0.2050	0.4000	0.1625	0.2364	
5	0.5249	0.5355	0.1741	0.3985	
6	0.2576	0.3784	0.1932	0.3063	
7	0.3682	0.2678	0.1556	0.3196	
8	0.5401	0.5390	0.2060	0.4378	
9	1.0580	1.1050	0.2210	0.4366	

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN
5	0.12440	NaN
6	0.08368	NaN
7	0.11510	NaN
8	0.10720	NaN
9	0.20750	NaN

[10 rows x 33 columns]

In [3]: *#Now we want to check how many columns and rows our data has, to do so we use the func*
df.shape

Out[3]: (569, 33)

In [4]: *#Now we will check for empty values in each column*

```
df.isna().sum()
```

```
Out[4]: id                                0
        diagnosis                         0
        radius_mean                      0
        texture_mean                     0
        perimeter_mean                   0
        area_mean                        0
        smoothness_mean                  0
        compactness_mean                 0
        concavity_mean                   0
        concave points_mean              0
        symmetry_mean                    0
        fractal_dimension_mean           0
        radius_se                        0
        texture_se                       0
        perimeter_se                     0
        area_se                          0
        smoothness_se                    0
        compactness_se                   0
        concavity_se                     0
        concave points_se                0
        symmetry_se                      0
        fractal_dimension_se             0
        radius_worst                     0
        texture_worst                    0
        perimeter_worst                  0
        area_worst                       0
        smoothness_worst                  0
        compactness_worst                 0
        concavity_worst                   0
        concave points_worst              0
        symmetry_worst                    0
        fractal_dimension_worst           0
        Unnamed: 32                      569
        dtype: int64
```

```
In [5]: #Since we have seen we have 569 and we find there is a column (Unnamed:32) with 569 mi.
        df.dropna(axis=1, inplace=True)
```

```
In [6]: #In order to prepare our data so machine learning algorithms "understand" it, we will
        df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
id                    569 non-null int64
diagnosis             569 non-null object
radius_mean           569 non-null float64
```

```

texture_mean          569 non-null float64
perimeter_mean        569 non-null float64
area_mean             569 non-null float64
smoothness_mean       569 non-null float64
compactness_mean      569 non-null float64
concavity_mean        569 non-null float64
concave points_mean   569 non-null float64
symmetry_mean         569 non-null float64
fractal_dimension_mean 569 non-null float64
radius_se            569 non-null float64
texture_se            569 non-null float64
perimeter_se         569 non-null float64
area_se              569 non-null float64
smoothness_se        569 non-null float64
compactness_se       569 non-null float64
concavity_se         569 non-null float64
concave points_se    569 non-null float64
symmetry_se          569 non-null float64
fractal_dimension_se  569 non-null float64
radius_worst         569 non-null float64
texture_worst        569 non-null float64
perimeter_worst      569 non-null float64
area_worst           569 non-null float64
smoothness_worst     569 non-null float64
compactness_worst    569 non-null float64
concavity_worst      569 non-null float64
concave points_worst 569 non-null float64
symmetry_worst       569 non-null float64
fractal_dimension_worst 569 non-null float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

```

```

In [7]: #Diagnosis contains M if tumour is Malign and B if Benign. We change these "object" va
        df['diagnosis'].replace(['M','B'],[1,0],inplace=True)

```

```

In [8]: #We observe values have changed
        df.diagnosis

```

```

Out[8]: 0      1
        1      1
        2      1
        3      1
        4      1
        ..
        564    1
        565    1
        566    1

```

```
567     1
568     0
Name: diagnosis, Length: 569, dtype: int64
```


1.0.5 5. Data Analysis and Visualization

After having cleaned and prepared our data so it is "ready" for our models to use, we are first going to analyze it in order to understand it better through plotting important information that will help us check how features behave and how they are correlated.

First of all, knowing our target variable is **diagnosis**, we are going to plot some information about it. Let's see which values this column contains and how many of them there are.

```
In [9]: #let's start visualizing the amount of Malignant and Benignant tumours in our dataset
sns.countplot(df['diagnosis'])
df['diagnosis'].value_counts()
```

```
Out[9]: 0     357
        1     212
Name: diagnosis, dtype: int64
```




As we see, approximately the 40% of tumours are detected as Malign while the 60% of them seem to be Bening. This might help us classify since both classes are almost equally represented.

Data Correlation Now let's look at how some features are related to having a benign or malign tumour. To do so, we are first going to plot a graph that will allow us to see how correlated features are with diagnosis.

```
In [10]: #Plot correlations with diagnosis column
corr = df.corr()['diagnosis'].sort_values(ascending=False)
corr.plot(kind='bar')
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1c6fc15ca90>
```



As we can see, almost every feature correlates positively with the diagnosis column and, these correlations are quite impressive, showing us there are quite some factors to take into account when deciding whether a tumour is benign or malign. That said we will do further investigations to those features whose correlation with diagnosis is greater than 0,6. In order to do so, we first need to check which features meet this requirement.


```
In [11]: #Plot only correlations greater than 0.6
corr[corr >= 0.6].plot(kind='bar')
print(corr[corr >= 0.6])
```

```
diagnosis          1.000000
concave points_worst 0.793566
perimeter_worst    0.782914
```

```

concave points_mean    0.776614
radius_worst           0.776454
perimeter_mean        0.742636
area_worst             0.733825
radius_mean            0.730029
area_mean              0.708984
concavity_mean         0.696360
concavity_worst        0.659610
Name: diagnosis, dtype: float64

```



Now that we know the most important features in this dataset, let's plot the correlation matrix to see how all features correlate with each other.

```


In [12]: #plot correlation matrix
plt.figure(figsize=(20,20))
sns.heatmap(df.iloc[:,1:32].corr(), annot=True, fmt='.0%')

```

```

Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1c688848518>

```

Data Standarization Having observed all the correlation between features, we are going to investigate a little bit more the ten features we selected before, these being "concave points_worst", "perimeter_worst", "concave points_mean", "radius_worst", "perimeter_mean", "area_worst", "radius_mean", "area_mean", "concavity_mean", "concavity_worst". However, before we do that, let's take a look at some of our data values.

```
In [13]: i_features = df.loc[:,["concave points_worst", "perimeter_worst", "concave points_mean",
                                "radius_worst", "perimeter_mean", "area_worst", "radius_mean", "area_mean",
                                "concavity_mean", "concavity_worst"]]
i_features.describe()
```

```
Out[13]:
```

	concave points_worst	perimeter_worst	concave points_mean	radius_worst	perimeter_mean	area_worst	radius_mean	area_mean	concavity_mean	concavity_worst
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000
mean	0.114606	107.261213	0.048919	0.014606	107.261213	0.014606	0.014606	0.014606	0.014606	0.014606

std	0.065732	33.602542	0.038803
min	0.000000	50.410000	0.000000
25%	0.064930	84.110000	0.020310
50%	0.099930	97.660000	0.033500
75%	0.161400	125.400000	0.074000
max	0.291000	251.200000	0.201200

	radius_worst	perimeter_mean	area_worst	radius_mean	area_mean	\
count	569.000000	569.000000	569.000000	569.000000	569.000000	
mean	16.269190	91.969033	880.583128	14.127292	654.889104	
std	4.833242	24.298981	569.356993	3.524049	351.914129	
min	7.930000	43.790000	185.200000	6.981000	143.500000	
25%	13.010000	75.170000	515.300000	11.700000	420.300000	
50%	14.970000	86.240000	686.500000	13.370000	551.100000	
75%	18.790000	104.100000	1084.000000	15.780000	782.700000	
max	36.040000	188.500000	4254.000000	28.110000	2501.000000	

	concavity_mean	concavity_worst
count	569.000000	569.000000
mean	0.088799	0.272188
std	0.079720	0.208624
min	0.000000	0.000000
25%	0.029560	0.114500
50%	0.061540	0.226700
75%	0.130700	0.382900
max	0.426800	1.252000

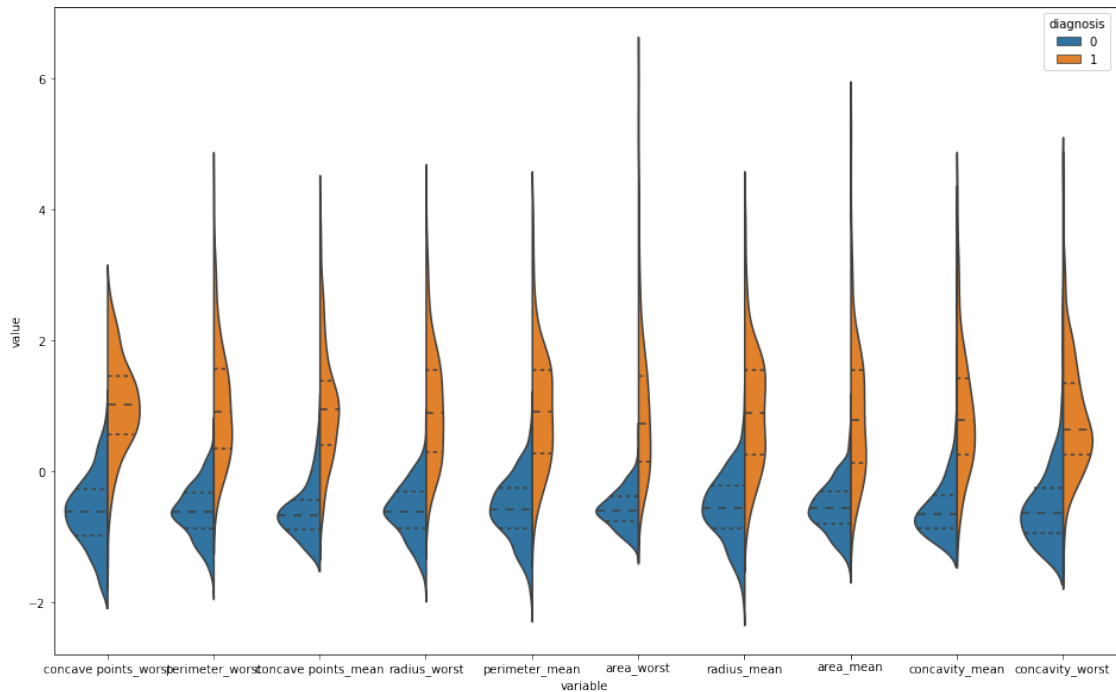
It can be observed by just looking at "max" values, concave points_worst max value is 0,291 while area_mean, for example, is 2501. That means if we want to group and show features together we really need to standarize or normalize our data so the difference between values isn't that big.

In this particular case, we choose to standarize data.

```
In [14]: #Data standarization and plotting information related to values
standarized_df = (i_features - i_features.mean()) / (i_features.std())
```

Data Plotting

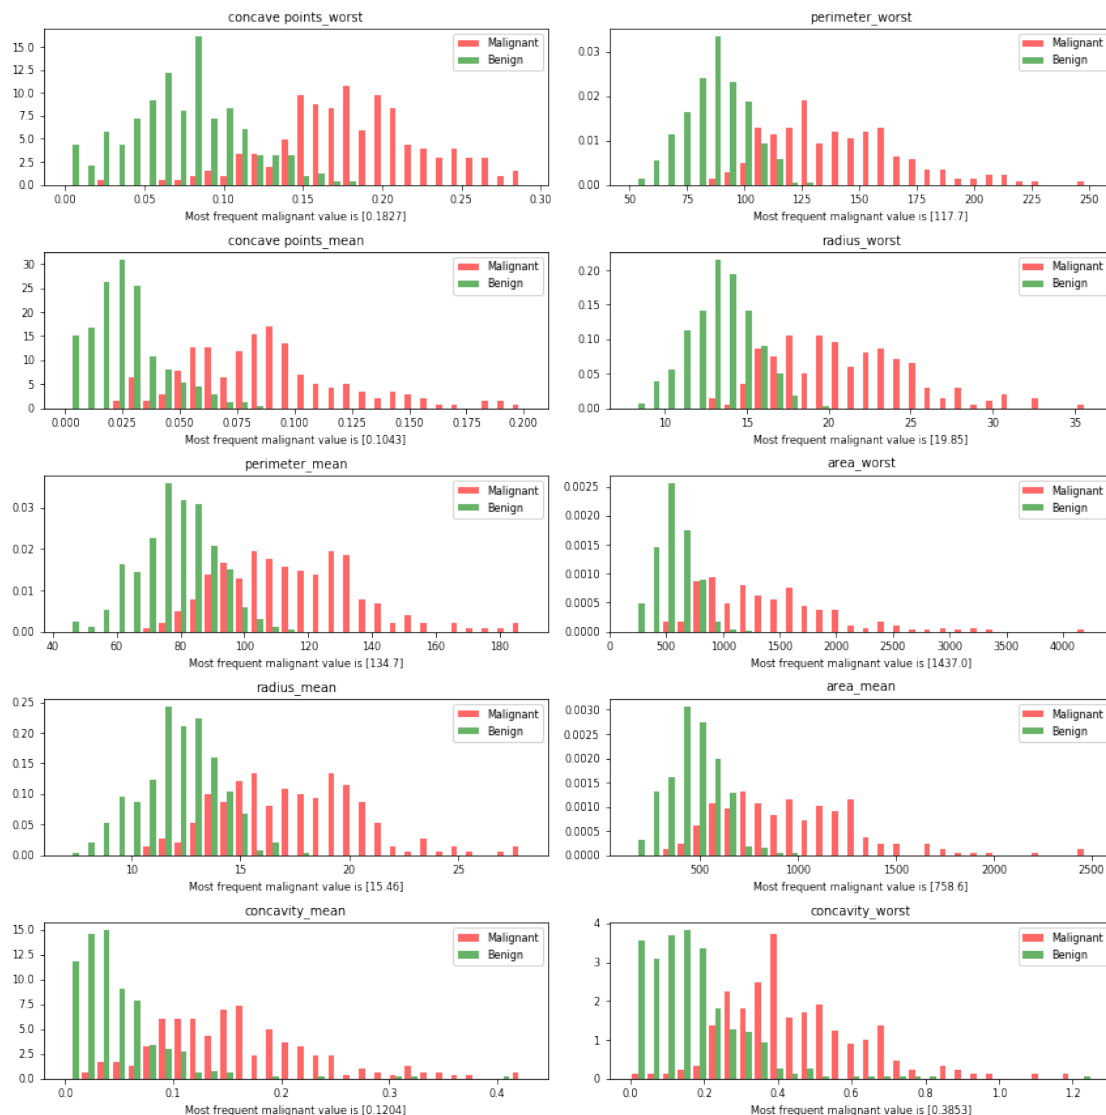
```
In [15]: #Violinplotting data
i_dia = df.diagnosis
standarized_df = pd.concat([i_dia,standarized_df],axis=1)
melted_data = pd.melt(standarized_df, id_vars = "diagnosis")
plt.figure(figsize = (16,10))
sns.violinplot(x = "variable", y = "value", hue="diagnosis", data = melted_data, split=True)
plt.show()
```



As we can see, it seems that the higher the values of these features the higher the chance the tumor is Malignant, that's how we can visually observe and confirm the positive correlation between these features and our diagnosis. Additionally, from the violinplot we can also see how "separated" are the benignant and malignant values, meaning these chosen features are a well representation of values distribution.

Now let's plot the histogram of these values so we can better observe their values and data distribution. In order to do so, we are going to separate, for each histogram, the values depending on the diagnosis column.

```
In [16]: #Here we are going to plot the histograms of the features we selected before and the
flistt = ["concave points_worst", "perimeter_worst", "concave points_mean", "radius_worst", "perimeter_mean", "area_worst", "radius_mean", "area_mean", "concavity_mean", "concavity_worst"]
malignant = df[df['diagnosis'] == 1]
benign = df[df['diagnosis'] == 0]
plt.rcParams.update({'font.size': 8})
plot, graphs = plt.subplots(nrows= 5, ncols= 2, figsize=(12,12))
graphs = graphs.flatten()
for i, graph in enumerate(graphs):
    graph.figure
    graph.hist([malignant[flistt[i]],benign[flistt[i]]], bins=30, normed = True, alpha=0.5)
    graph.legend(loc='upper right')
    graph.set_xlabel("Most frequent malignant value is " + str(malignant[flistt[i]].value))
    graph.set_title(flistt[i])
plt.tight_layout()
```



From these ten graphs we can observe, as we said before, these features might be useful in predicting whether a patient has cancer or not due to the distinct grouping between malignant and benign. We can also see the most frequent malignant value of these features under each graph, although visually they may vary since we might not have an optimal bin width, but most of them match with the values histograms show.

1.0.6 6. Feature selection

Feature selection is the process where we automatically or manually select those features that really help when making predictions, since they have a decent correlation with the target variable. It is one of the core concepts in machine learning since it really impacts the performance of our model and can also prevent overfitting, improve accuracy and reduce training time. For example, on this problem feature selection might very important since having a machine learning algorithm with 200 features and an accuracy of 97% or having another algorithm with 10 features and an

accuracy of 96%, it is pretty clear we should go for the second one, since being able to detect whether a tumor is benign or malignant with only a few features instead of way more can save us a pretty decent amount of time and ease things.

In order to do this feature selection we are going to use two different methods. In the first method we are going to manually select the features we want to train our models with, which will be the same ten features we selected before, and in the second method we will automatically select them by using recursive feature elimination from random forest.

We will use three different algorithms to train our models with our manually selected data and, after that, we are going to select the best one and compare its results with the results obtained with automatic selection.

1.0.7 7. Classification

1.0.8 Manually Selected Features

Before doing any kind of training we need to select the features we want to train our algorithms with and separate them from the target variable. In order to do so we are going to split our data into training and set using the function `train_test_split`.

```
In [17]: #Data separation
```

```
y = df.iloc[:,1]
```

```
x = df.loc[:,["concave points_worst", "perimeter_worst", "concave points_mean", "radiation_worst"]]
```

```
In [18]: #splitting data
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=42)
```

Support Vector Machine

```
In [19]: #We will train our model with different parameters and the best "combination" will be selected
```

```
parameters = {'kernel':('linear', 'rbf'), 'C':[1, 10]}
```

```
svc = svm.SVC(gamma="scale")
```

```
clf = GridSearchCV(svc, parameters, cv=5)
```

```
clf.fit(x_train, y_train)
```

```
print(clf.best_estimator_)
```

```
SVC(C=10, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='linear',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
In [20]: #Print some metrics in a table
```

```
ptbl = PrettyTable()
```

```
ptbl.field_names = ["Accuracy", "Recall", "F1Score"]
```

```
ptbl.add_row([accuracy_score(y_test, clf.predict(x_test)),
```

```
              recall_score(y_test, clf.predict(x_test)), f1_score(y_test, clf.predict(x_test))])
```

```
print(ptbl)
```

Accuracy	Recall	F1Score
0.9707602339181286	0.9682539682539683	0.9606299212598425

Gaussian Naive Bayes

```
In [21]: g_nb = GaussianNB()
         g_nb.fit(x_train, y_train)
```

```
Out[21]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [22]: #Print some metrics in a table
         ptbl = PrettyTable()
         ptbl.field_names = ["Accuracy", "Recall", "F1Score"]
         ptbl.add_row([accuracy_score(y_test, g_nb.predict(x_test)),
                        recall_score(y_test, g_nb.predict(x_test)), f1_score(y_test, g_nb.predict(x_test))])
         print(ptbl)
```

Accuracy	Recall	F1Score
0.9590643274853801	0.9206349206349206	0.943089430894309

Decision Tree

```
In [23]: #Decision tree using entropy
         decision_tree = DecisionTreeClassifier(criterion = 'entropy', random_state = 42)
         decision_tree.fit(x_train, y_train)
```

```
Out[23]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=42, splitter='best')
```

```
In [24]: #Print some metrics in a table
         ptbl = PrettyTable()
         ptbl.field_names = ["Accuracy", "Recall", "F1Score"]
         ptbl.add_row([accuracy_score(y_test, decision_tree.predict(x_test)),
                        recall_score(y_test, decision_tree.predict(x_test)), f1_score(y_test, decision_tree.predict(x_test))])
         print(ptbl)
```

Accuracy	Recall	F1Score
0.935672514619883	0.9206349206349206	0.9133858267716536

1.0.9 Automatically Selected Features

Having trained our models with manually selected data we obtained some good results. Now we are going to follow the same steps but using recursive feature elimination with cross validation and random forest in order to check how many of these features we do really need in order to have the optimal accuracy in our model. In order to do so, we will use our original data with all features and we're going to split it again.

```
In [25]: yrf = df.iloc[:,1]
        xrf = df.iloc[:,2:31]
        xrf_train, xrf_test, yrf_train, yrf_test = train_test_split(xrf, yrf, test_size=0.3, r
```

```
In [26]: random_forest = RandomForestClassifier()
        rfecv = RFECV(estimator=random_forest, step=1, cv=5, scoring='accuracy')
        rfecv = rfecv.fit(xrf_train, yrf_train)
        print('Total features selected:', rfecv.n_features_)
        print('Selected features:', xrf_train.columns[rfecv.support_])
```

Total features selected: 22

Selected features: Index(['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave points_mean', 'fractal_dimension_mean', 'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'concavity_se', 'radius_worst', 'texture_worst', 'perimeter_worst', 'area_worst', 'smoothness_worst', 'concavity_worst', 'concave points_worst', 'symmetry_worst'], dtype='object')

```
In [27]: #Print some metrics in a table
        ptbl = PrettyTable()
        ptbl.field_names = ["Accuracy", "Recall", "F1Score"]
        ptbl.add_row([accuracy_score(yrf_test, rfecv.predict(xrf_test)),
                        recall_score(yrf_test, rfecv.predict(xrf_test)), f1_score(yrf_test, rfecv.p
        print(ptbl)
```

Accuracy	Recall	F1Score
0.9590643274853801	0.9206349206349206	0.943089430894309


1.0.10 8. Results

Apart from the scores we have shown in the last section, it's time that we make predictions using our models. In order to do so, we are going to select the support vector machine model (since it's the one, in manual feature selection, we obtained the best results) and compare it with our random forest that automatically selects the "optimal" features.

These predictions will be plotted as confusion matrices showing the relation between predicted and true values, that way we can clearly and visually observe True/False Positives and Negatives.

Support Vector Machine

```
In [28]: #Confussion Matrix for the SVM
label_aux = plt.subplot()
svm_prediction = clf.predict(x_test)
cm_svm = confusion_matrix(y_test,svm_prediction)
cm_svm_m = pd.DataFrame(cm_svm, index = ['B','M'], columns = ['B','M'])
sns.heatmap(cm_svm_m,annot=True,fmt="d", cbar=False)
label_aux.set_xlabel('Predicted Value');label_aux.set_ylabel('True Value');
```




Random Forest

```
In [29]: #Confussion Matrix for the Random Forest
label_aux = plt.subplot()
rf_prediction = rfecv.predict(xrf_test)
cm_rf = confusion_matrix(yrf_test,rf_prediction)
```



```
cm_rf_m = pd.DataFrame(cm_rf, index = ['B','M'], columns = ['B','M'])
sns.heatmap(cm_rf_m,annot=True,fmt="d", cbar=False)
label_aux.set_xlabel('Predicted Value');label_aux.set_ylabel('True Value');
```



1.0.11 9. Conclusions

After having obtained all the results through our models and plots, these are some things we can say about this problem and solution:

- Support Vector Machine and Random Forest classifiers seem to be the best ones in this dataset when trying to decide whether a tumour is benign or malign.
- A good amount of columns are redundant and do not give enough information at all since their correlation with diagnosis is close to none and having them in our dataset does not improve results.
- Manual feature selection can lead us to obtaining very good results if done properly, sometimes these results are even better than the ones obtained by automatical selection.
- When looking at tumour features that belong to "mean" and "worst" the ones that might give the most valuable information seem to be, especially: radius, area, perimeter, concave points and concavity.
- Around 2/3 of detected tumours are benign, so early detection can increase survival chances.
- In order to slightly improve these results, there seems to be a kernel in Kaggle where Spearman's Rank Correlation is used instead of Pearson's (provided by the heatmap) and seems to find stronger correlations since the relationship between variables isn't linear at all. However, the results barely change.

If someone wants to know a little bit more of this problem, the following link will take you the most voted kernel on Kaggle.